



ELG5166

Spark APIs



NOVEMBER 3, 2022

ALI AMIN EL-SAYED MAHMOUD EL-SHERIF
300327246

Question 1:

```
1  /* Q) 1 */
2  val sales = spark
3  .read
4  .option("inferSchema", "true")
5  .option("header", "true")
6  .csv("/FileStore/tables/Data/Sales.csv")
7  sales.count()
8  sales.createOrReplaceTempView("Sales")
9
10 sales.show(3)
```

▸ (5) Spark Jobs

▮ sales: org.apache.spark.sql.DataFrame = [Transaction_date: string, Product: string ... 10 more fields]

```
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
--+
|Transaction_date| Product|Price|Payment_Type|      Name|      City|  State|      Country|Account_Created|  Last_Login|Latitude| Longitu
de|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
--+
|      1/2/09 6:17|Product1| 1200|  Mastercard|    carolina|    Basildon|England|United Kingdom|  1/2/09 6:00|  1/2/09 6:08|  51.5|-1.11666
67|
|      1/2/09 4:53|Product1| 1200|      Visa|    Betina|Parkville|    ...|    MO| United States|  1/2/09 4:42|  1/2/09 7:49|  39.195|-94.681
94|
|      1/2/09 13:08|Product1| 1200|  Mastercard|Federica e Andrea|Astoria|    ...|    OR| United States|  1/1/09 16:21|1/3/09 12:32|46.18806| -123.
83|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
--+
only showing top 3 rows
```

```
1  val flightData = spark
2  .read
3  .option("inferSchema", "true")
4  .option("header", "true")
5  .csv("/FileStore/tables/flight_data.csv")
6  .coalesce(5)
7  flightData.count()
8  flightData.cache()
9  flightData.createOrReplaceTempView("flight")
10
```

▸ (4) Spark Jobs

▮ flightData: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [DEST_COUNTRY_NAME: string, ORIGIN_COUNTRY_NAME: string ... 1 more field]

flightData: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [DEST_COUNTRY_NAME: string, ORIGIN_COUNTRY_NAME: string ... 1 more field]

Cmd 3

```
1  /*Q)1 b-a */
2  import org.apache.spark.sql.functions.min
3  var query = spark.sql("select min(count) from flight")
4  query.show()
```

▸ (2) Spark Jobs

▮ query: org.apache.spark.sql.DataFrame = [min(count): Integer]

```
+-----+
|min(count)|
+-----+
|          1|
+-----+
```

```
import org.apache.spark.sql.functions.min
query: org.apache.spark.sql.DataFrame = [min(count): int]
```

Cmd 4

```
1  %sql
2  /*Q)1 b-b */
3  select min(count) from flight
```

▸ (2) Spark Jobs

min(count)
1

Showing 1 row. | 0.56 seconds runtime

Refreshed now

Question 2:

Cmd 5

```
1 %python
2 #Q)2
3 flight1 = spark.read.option("inferSchema", "true").option("header", "true").csv("/FileStore/tables/flight_data.csv")
4 flight1.createOrReplaceTempView("flight_data_2015")
```

▶ (7) Spark Inhe

Cmd 6

```
1 %sql
2 /*Q)2 a-a */
3 select sum(count) as total_number_flights,ORIGIN_COUNTRY_NAME from flight_data_2015 group by ORIGIN_COUNTRY_NAME order by total_number_flights limit 3
```

▶ (2) Spark Jobs

Table +

	total_number_flights	ORIGIN_COUNTRY_NAME
1	1	Singapore
2	1	Lithuania
3	1	Croatia

Showing all 3 rows. | 2.31 seconds runtime Refreshed now

Cmd 7

```
1 /*Q)2 a-b */
2 var query = spark.sql("select sum(count) as total_number_flights,ORIGIN_COUNTRY_NAME from flight_data_2015 group by ORIGIN_COUNTRY_NAME order by
3 total_number_flights limit 3")
4 query.show()
```

▶ (2) Spark Jobs

query: org.apache.spark.sql.DataFrame = [total_number_flights: long, ORIGIN_COUNTRY_NAME: string]

```
+-----+-----+
|total_number_flights|ORIGIN_COUNTRY_NAME|
+-----+-----+
|1|Singapore|
|1|Lithuania|
|1|Croatia|
+-----+-----+
```

query: org.apache.spark.sql.DataFrame = [total_number_flights: bigint, ORIGIN_COUNTRY_NAME: string]

```
1 /*Q)2 a-c */
2 var query = spark.sql("select sum(count) as total_number_flights,ORIGIN_COUNTRY_NAME from flight_data_2015 group by ORIGIN_COUNTRY_NAME order by
3 total_number_flights limit 3")
4 query.show()
5
6 flightData
7   .groupBy("ORIGIN_COUNTRY_NAME")
8   .sum("count")
9   .sort("total_number_flights")
10  .limit(3)
11
12
```

▶ (2) Spark Jobs

```
+-----+-----+
|total_number_flights|ORIGIN_COUNTRY_NAME|
+-----+-----+
|1|Singapore|
|1|Lithuania|
|1|Croatia|
+-----+-----+
```

⊞ AnalysisException: Column 'total_number_flights' does not exist. Did you mean one of the following? [sum(count), ORIGIN_COUNTRY_NAME];
'Sort ['total_number_flights ASC NULLS FIRST], true
+- Aggregate [ORIGIN_COUNTRY_NAME#192], [ORIGIN_COUNTRY_NAME#192, sum(count#193) AS sum(count)#503L]
+- Repartition 5, false
+- Relation [DEST_COUNTRY_NAME#191,ORIGIN_COUNTRY_NAME#192,count#193] csv

```

1  /*Q)2 b */
2  var query = spark.sql("select sum(count) as total_number_flights,ORIGIN_COUNTRY_NAME from flight_data_2015 group by ORIGIN_COUNTRY_NAME order by
3  total_number_flights limit 3")
4  query.show()
5  query.explain()

```

► (2) Spark Jobs

► query: org.apache.spark.sql.DataFrame = [total_number_flights: long, ORIGIN_COUNTRY_NAME: string]

total_number_flights	ORIGIN_COUNTRY_NAME
1	Singapore
1	Lithuania
1	Croatia

== Physical Plan ==

AdaptiveSparkPlan isFinalPlan=false

+-- TakeOrderedAndProject(limit=3, orderBy=[total_number_flights#507L ASC NULLS FIRST], output=[total_number_flights#507L,ORIGIN_COUNTRY_NAME#421])

+-- HashAggregate(keys=[ORIGIN_COUNTRY_NAME#421], functions=[finalmerge_sum(merge sum#518L) AS sum(count#422)#508L])

+-- Exchange hashpartitioning(ORIGIN_COUNTRY_NAME#421, 200), ENSURE_REQUIREMENTS, [id=#792]

+-- HashAggregate(keys=[ORIGIN_COUNTRY_NAME#421], functions=[partial_sum(count#422) AS sum#518L])

+-- FileScan csv [ORIGIN_COUNTRY_NAME#421,count#422] Batched: false, DataFilters: [], Format: CSV, Location: InMemoryFileIndex(1 paths)[dbfs:/FileStore/tables/flight_data.csv], PartitionFilters: [], PushedFilters: [], ReadSchema: struct<ORIGIN_COUNTRY_NAME:string,count:int>

query: org.apache.spark.sql.DataFrame = [total_number_flights: bigint, ORIGIN_COUNTRY_NAME: string]

Question 3:

- (2) Spark Jobs
- flightsDataFile: org.apache.spark.sql.DataFrame = [DEST_COUNTRY_NAME: string, ORIGIN_COUNTRY_NAME: string ... 1 more field]
- flights: org.apache.spark.sql.Dataset[Flight] = [DEST_COUNTRY_NAME: string, ORIGIN_COUNTRY_NAME: string ... 1 more field]
- flights_Metadata: org.apache.spark.sql.Dataset[FlightMetadata] = [count: long, random: long]

count	random
0	-5790796081456756646
1	-7406709754318658195
2	3723754083696420047
3	-3678622977270750094
4	-8820342020588039680
5	554201728505525082
6	-4210117342112615418
7	8640157511789465638
8	5967035757421660660
9	3968210518744556750
10	3652462764861634911
11	4179064758276548709
12	3928348356421848608
13	-4441871604060770695
14	-5151248049971105193
15	-1522132841203857217
16	4844185957588460211
17	-5312010116759265553

```

1  /*Q)3 a*/
2
3  case class Flight(DEST_COUNTRY_NAME: String,
4                    ORIGIN_COUNTRY_NAME: String, count: BigInt)
5
6  /* Reads the Parquet file */
7  val flightsDataFile = spark.read
8    .parquet("/FileStore/tables/Data/flight-data/parquet/2010-summary.parquet/part_r_00000_1a9822ba_b8fb_4d8e_844a_ea30d0801b9e_gz.parquet")
9
10 /* map the file to a typed dataset using Flight class*/
11 val flights = flightsDataFile.as[Flight]
12
13 /* create a structured dataset of the Type FlightMetadata */
14 case class FlightMetadata(count: BigInt, random: BigInt)
15
16 /* 500 entries with an index column and another with random numbers */
17 val flights_Metadata = spark.range(500)
18   .map(x => (x, scala.util.Random.nextLong))
19   .withColumnRenamed("_1", "count")
20   .withColumnRenamed("_2", "random")
21   .as[FlightMetadata]
22 flights_Metadata.show()
23

```

```

1  /*Q)3 b*/
2  val flights_2 = flights
3    .joinWith(flights_Metadata, flights.col("count") === flights_Metadata.col("count"))
4    .withColumnRenamed("_1", "count")
5    .withColumnRenamed("_2", "random")
6    flights_2.show()
7

```

flights_2: org.apache.spark.sql.DataFrame = [count: struct, random: struct]

	count	random
{United States, U...}	{1, -208177811221...}	
{United States, F...}	{1, -208177811221...}	
{Bulgaria, United...}	{1, -208177811221...}	
{United States, S...}	{1, -208177811221...}	
{United States, C...}	{1, -208177811221...}	
{United States, B...}	{1, -208177811221...}	
{United States, I...}	{1, -208177811221...}	
{Vietnam, United ...}	{1, -208177811221...}	
{United States, P...}	{1, -208177811221...}	
{United States, B...}	{1, -208177811221...}	
{United States, G...}	{1, -208177811221...}	
{United States, L...}	{1, -208177811221...}	
{Malaysia, United...}	{1, -208177811221...}	
{United States, S...}	{1, -208177811221...}	
{United States, A...}	{1, -208177811221...}	
{United States, E...}	{1, -208177811221...}	
{United States, V...}	{1, -208177811221...}	
{Liberia, United ...}	{1, -208177811221...}	

Question 4:

```

1  %sql
2  /* Q)4 a */
3
4  DROP TABLE IF EXISTS flights;
5  CREATE TABLE flights (
6    DEST_COUNTRY_NAME STRING, ORIGIN_COUNTRY_NAME STRING, count LONG)
7  USING JSON OPTIONS (path '/FileStore/tables/Data/flight-data/json/*.json')

```

OK

```

1  /* Q)4 b*/
2
3  flights.rdd.getNumPartitions

```

res18: Int = 1

Question 5:

```

1  /* Q)5 a*/
2
3  val textFile = spark.sparkContext.textFile("/FileStore/tables/Data/Adult.csv")
4  textFile.count()

```

► (1) Spark Jobs

textFile: org.apache.spark.rdd.RDD[String] = /FileStore/tables/Data/Adult.csv MapPartitionsRDD[108] at textFile at command-3522973057872627:3
res19: Long = 32562

```

1  /* Q)5 b*/
2
3  val count = textFile.flatMap(line => line.split(","))
4                      .map(word => (word, 1))
5                      .reduceByKey(_+_).takeOrdered(5)(Ordering[Int].on(_._2))
6  count.foreach(println)
7

```

► (1) Spark Jobs

(87,1)
(marital_status,1)
(age,1)
(Holand-Nether,1)
(occupation,1)
count: Array[(String, Int)] = Array((87,1), (marital_status,1), (age,1), (Holand-Nether,1), (occupation,1))

Question 6:

```

1  /* Q)6 */
2
3  val dataframe = spark.read.format("csv")
4                  .option("header", "true")
5                  .option("inferSchema", "true")
6                  .load("/FileStore/tables/2010_12_10.csv")
7
8  dataframe.printSchema()
9
10 // Create a view.
11 dataframe.createOrReplaceTempView("dfTable")

```

► (2) Spark Jobs

dataframe: org.apache.spark.sql.DataFrame = [InvoiceNo: string, StockCode: string ... 6 more fields]

```

root
 |-- InvoiceNo: string (nullable = true)
 |-- StockCode: string (nullable = true)
 |-- Description: string (nullable = true)
 |-- Quantity: integer (nullable = true)
 |-- InvoiceDate: timestamp (nullable = true)
 |-- UnitPrice: double (nullable = true)
 |-- CustomerID: double (nullable = true)
 |-- Country: string (nullable = true)

dataframe: org.apache.spark.sql.DataFrame = [InvoiceNo: string, StockCode: string ... 6 more fields]

```

```
Cmd 17
1 /*Q)6 a*/
2
3 import org.apache.spark.sql.functions.col
4 dataframe.where(col("Quantity").equalTo(12) && col("UnitPrice").gt(2))
5 .select("*")
6 .show(5, false)
7
8
```

► (1) Spark Jobs

InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
538172	22041	"RECORD FRAME 7"" SINGLE SIZE "	12	2010-12-10 09:33:00	2.55	15805.0	United Kingdom
538172	84558A	3D DOG PICTURE PLAYING CARDS	12	2010-12-10 09:33:00	2.95	15805.0	United Kingdom
538174	22326	ROUND SNACK BOXES SET OF4 WOODLAND	12	2010-12-10 09:35:00	2.95	12471.0	Germany
538174	22472	TV DINNER TRAY DOLLY GIRL	12	2010-12-10 09:35:00	4.95	12471.0	Germany
538174	22445	PENCIL CASE LIFE IS BEAUTIFUL	12	2010-12-10 09:35:00	2.95	12471.0	Germany

only showing top 5 rows

import org.apache.spark.sql.functions.col

```
Cmd 18
1 /*Q)6 b*/
2
3 import org.apache.spark.sql.functions.split
4 dataframe.select(split(col("Description"), " ").alias("Detailed description")).show(false)
5
```

► (1) Spark Jobs

Detailed description
[HAWAIIAN, GRASS, SKIRT,]
[CHILLI, LIGHTS]
["RECORD, FRAME, 7"" SINGLE, SIZE, "]
[3D, DOG, PICTURE, PLAYING, CARDS]
[60, CAKE, CASES, VINTAGE, CHRISTMAS]
[PAPER, CHAIN, KIT, VINTAGE, CHRISTMAS]
[CHRISTMAS, TOILET, ROLL]
["ASSORTED, FLOWER, COLOUR, ""LEIS"""]
[I, CAN, ONLY, PLEASE, ONE, PERSON, MUG]
[HAND, WARMER, BABUSHKA, DESIGN]
[BLUE, HARMONICA, IN, BOX,]
null
[SET, OF, 72, RETROSPOT, PAPER, , DOILIES]
[PACK, OF, 72, RETROSPOT, CAKE, CASES]
[WOODLAND, DESIGN, , COTTON, TOTE, BAG]
[BIG, DOUGHNUT, FRIDGE, MAGNETS]
[RECYCLED, PENCIL, WITH, RABBIT, ERASER]

```
Cmd 19
1 /*Q)6 2-a*/
2
```

```
1 %fs
2 ls /FileStore/tables
```

Table +				
	path	name	size	modificationTime
1	dbfs:/FileStore/tables/2010_12_01.csv	2010_12_01.csv	275001	1667513508000
2	dbfs:/FileStore/tables/2010_12_02.csv	2010_12_02.csv	191826	1667513508000
3	dbfs:/FileStore/tables/2010_12_10.csv	2010_12_10.csv	241468	1667513718000
4	dbfs:/FileStore/tables/Data/	Data/	0	0
5	dbfs:/FileStore/tables/flight_data.csv	flight_data.csv	7080	1666034682000

Showing all 5 rows. | 1.13 seconds runtime

Refreshed 2 minutes ago

```

Cmd 21
1  /*Q)6 2-b*/
2
3  val dataset = spark.read.format("csv")
4    .option("header", "true")
5    .option("inferSchema", "true")
6    .load("/FileStore/tables/2010_12_02.csv")
7
8  dataset.createOrReplaceTempView("dfTable")
9
10 import org.apache.spark.sql.functions.{sum,min, max,avg}
11
12 dataset.select(sum("UnitPrice"),min("UnitPrice"), max("UnitPrice"),avg("UnitPrice")).show()

```

▶ (4) Spark Jobs

dataset: org.apache.spark.sql.DataFrame = [InvoiceNo: string, StockCode: string ... 6 more fields]

sum(UnitPrice)	min(UnitPrice)	max(UnitPrice)	avg(UnitPrice)
6815.679999999958	0.0	295.0	3.231711711711692

Question 7:

```

Cmd 22
1  /* Q)7 */
2
3  val dataset = spark.read.format("csv")
4    .option("header", "true")
5    .option("inferSchema", "true")
6    .load("/FileStore/tables/2010_12_01.csv")
7
8  dataset.createOrReplaceTempView("RetailView")
9  dataset.count()

```

▶ (4) Spark Jobs

dataset: org.apache.spark.sql.DataFrame = [InvoiceNo: string, StockCode: string ... 6 more fields]

res29: Long = 3108

```

Cmd 23
1  /*Q)7 a*/
2  val datadescription= dataset.filter(x =>
3  {
4      val allStr = x.getString(2).split(" ")
5      var check = true
6      for(x <- allStr ){
7          if (!x(0).isUpper)
8          {
9              check=false
10             }
11         }
12     }
13     (check)
14 })
15
16 datadescription.show(5,false)

```

▶ (1) Spark Jobs

datadescription: org.apache.spark.sql.Dataset[org.apache.spark.sql.Row] = [InvoiceNo: string, StockCode: string ... 6 more fields]

InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	2010-12-01 08:26:00	2.55	17850.0	United Kingdom
536365	71053	WHITE METAL LANTERN	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom
536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	2010-12-01 08:26:00	2.75	17850.0	United Kingdom
536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom
536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	2010-12-01 08:26:00	3.39	17850.0	United Kingdom

only showing top 5 rows

Cmd 24

```
1 /*Q)7 b*/  
2 /*Drop a any row that contains NULLs using any*/  
3  
4 dataset.na.drop("any")  
5
```

▶ res33: org.apache.spark.sql.DataFrame = [InvoiceNo: string, StockCode: string ... 6 more fields]
res33: org.apache.spark.sql.DataFrame = [InvoiceNo: string, StockCode: string ... 6 more fields]

```
1 /*Drop a any row that contains NULLs using all*/  
2  
3 dataset.na.drop("all")
```

▶ res35: org.apache.spark.sql.DataFrame = [InvoiceNo: string, StockCode: string ... 6 more fields]
res35: org.apache.spark.sql.DataFrame = [InvoiceNo: string, StockCode: string ... 6 more fields]