




uOttawa



ELG5166 Cloud Analytics

Take Home Exam



Instructor: Dr. Benjamin Eze

Student Name: Ali El-Sherif

ID: 300327246

Table of Contents

Personal Ethics & Academic Integrity Statement.....	2
Question 1:.....	3
First assumption (All nodes will fail at the same time)	3
Second assumption (All nodes failure calculated throughout the year).....	3
Question 2:.....	3
Question 3:.....	5
a).....	5
b).....	5
Question 4:.....	6
a).....	6
b).....	6
c).....	6
d).....	6
e).....	7
References.....	8

Personal Ethics & Academic Integrity Statement

By typing in my name and student ID on this form and submitting it electronically, I am attesting to the fact that I have reviewed not only my work but the work of my team member, in its entirety.

I attest to the fact that my work in this project adheres to the fraud policies as outlined in the Academic Regulations in the University's Graduate Studies Calendar. I further attest that I have knowledge of and have respected the "Beware of Plagiarism" brochure for the university. To the best of my knowledge, I also believe that each of my group colleagues has also met the aforementioned requirements and regulations. I understand that if my group assignment is submitted without a completed copy of this Personal Work Statement from each group member, it will be interpreted by the school that the missing student(s) name is confirmation of non-participation of the aforementioned student(s) in the required work. We, by typing in our names and student IDs on this form and submitting it electronically,

- warrant that the work submitted herein is our own group members' work and not the work of others.
- acknowledge that we have read and understood the University Regulations on Academic Misconduct.
- acknowledge that it is a breach of University Regulations to give or receive unauthorized and/or unacknowledged assistance on a graded piece of work.

Question 1:

First assumption (All nodes will fail at the same time)

$$\text{Probability} = \left(\frac{24}{365}\right)^5 = 1.229 \times 10^{-6}$$

$$\text{Uptime} = 1 - \text{probability} = 1 - 1.229 \times 10^{-6} = 0.9999987$$

$$\text{Uptime percentage} = \text{uptime} \times 100 = 99.99987\%$$

Availability = 5 nines

Second assumption (All nodes failure calculated throughout the year)

$$\text{Probability} = \left(\frac{24 \times 5}{365 \times 5}\right) = 0.065753$$

$$\text{Uptime} = 1 - \text{probability} = 1 - 0.065753 = 93.4247$$

$$\text{Uptime percentage} = \text{uptime} \times 100 = 93.4247\%$$

Availability = one nine

Question 2:

Reading the data as spark dataframe:

```
val main_data = spark.read.option("inferSchema","true").option("header","true").csv ("/FileStore/tables/retail-data/daily/*.csv")
display (main_data)
```

▶ (3) Spark Jobs

▶ main_data: org.apache.spark.sql.DataFrame = [InvoiceNo: string, StockCode: string ... 6 more fields]

Table ▾ +

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
1	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	12/1/2010 8:26	2.55	17850	United Kingdom
2	536365	71053	WHITE METAL LANTERN	6	12/1/2010 8:26	3.39	17850	United Kingdom
3	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	12/1/2010 8:26	2.75	17850	United Kingdom
4	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	12/1/2010 8:26	3.39	17850	United Kingdom
5	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	12/1/2010 8:26	3.39	17850	United Kingdom
6	536365	22752	SET 7 BABUSHKA NESTING BOXES	2	12/1/2010 8:26	7.65	17850	United Kingdom
7	536365	21730	GLASS STAR FROSTED T-LIGHT HOLDER	6	12/1/2010 8:26	4.25	17850	United Kinadom

Truncated results, showing first 1,000 rows. | 7.37 seconds runtime

Refreshed 2 minutes ago

Command took 7.37 seconds -- by alielsherifeng@gmail.com at 12/16/2022, 11:20:53 PM on Ali El-Sherif

Change date column from string to date format:

```
import spark.sqlContext.implicits._
import org.apache.spark.sql.functions._
spark.conf.set("spark.sql.legacy.timeParserPolicy", "LEGACY")

import spark.sqlContext.implicits._
import org.apache.spark.sql.functions._

Command took 2.02 seconds -- by alielsherifeng@gmail.com at 12/16/2022, 11:20:53 PM on Ali El-Sherif
```

Cmd 4

```
val dated_data = main_data.withColumn("date", to_date(col("InvoiceDate"), "MM/dd/yyyy"))
display(dated_data)
```

▶ (1) Spark Jobs

▶  dated_data: org.apache.spark.sql.DataFrame = [InvoiceNo: string, StockCode: string ... 7 more fields]

Table +

	InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country
1	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	12/1/2010 8:26	2.55	17850	United Kingdom
2	536365	71053	WHITE METAL LANTERN	6	12/1/2010 8:26	3.39	17850	United Kingdom
3	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	12/1/2010 8:26	2.75	17850	United Kingdom
4	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	12/1/2010 8:26	3.39	17850	United Kingdom
5	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	12/1/2010 8:26	3.39	17850	United Kingdom
6	536365	22752	SET 7 BABUSHKA NESTING BOXES	2	12/1/2010 8:26	7.65	17850	United Kingdom
7	536365	21730	GLASS STAR FROSTED T-LIGHT HOLDER	6	12/1/2010 8:26	4.25	17850	United Kinadom

↓

▼

Truncated results, showing first 1,000 rows. | 1.20 seconds runtime

Refreshed 2 minutes ago

Command took 1.20 seconds -- by alielsherifeng@gmail.com at 12/16/2022, 11:20:53 PM on Ali El-Sherif

Add a month column as follows:

```
val sorted_data = dated_data.withColumn("month", month(col("date")))
sorted_data.show()
```

▶ (1) Spark Jobs

▶  sorted_data: org.apache.spark.sql.DataFrame = [InvoiceNo: string, StockCode: string ... 8 more fields]

InvoiceNo	StockCode	Description	Quantity	InvoiceDate	UnitPrice	CustomerID	Country	date	month
536365	85123A	WHITE HANGING HEA...	6	12/1/2010 8:26	2.55	17850	United Kingdom	2010-12-01	12
536365	71053	WHITE METAL LANTERN	6	12/1/2010 8:26	3.39	17850	United Kingdom	2010-12-01	12
536365	84406B	CREAM CUPID HEART...	8	12/1/2010 8:26	2.75	17850	United Kingdom	2010-12-01	12
536365	84029G	KNITTED UNION FLA...	6	12/1/2010 8:26	3.39	17850	United Kingdom	2010-12-01	12
536365	84029E	RED WOOLLY HOTTIE...	6	12/1/2010 8:26	3.39	17850	United Kingdom	2010-12-01	12
536365	22752	SET 7 BABUSHKA NE...	2	12/1/2010 8:26	7.65	17850	United Kingdom	2010-12-01	12
536365	21730	GLASS STAR FROSTE...	6	12/1/2010 8:26	4.25	17850	United Kingdom	2010-12-01	12
536366	22633	HAND WARMER UNION...	6	12/1/2010 8:28	1.85	17850	United Kingdom	2010-12-01	12
536366	22632	HAND WARMER RED P...	6	12/1/2010 8:28	1.85	17850	United Kingdom	2010-12-01	12
536367	84879	ASSORTED COLOUR B...	32	12/1/2010 8:34	1.69	13047	United Kingdom	2010-12-01	12
536367	22745	POPPY'S PLAYHOUSE...	6	12/1/2010 8:34	2.1	13047	United Kingdom	2010-12-01	12
536367	22748	POPPY'S PLAYHOUSE...	6	12/1/2010 8:34	2.1	13047	United Kingdom	2010-12-01	12
536367	22749	FELTCRAFT PRINCES...	8	12/1/2010 8:34	3.75	13047	United Kingdom	2010-12-01	12
536367	22310	IVORY KNITTED MUG...	6	12/1/2010 8:34	1.65	13047	United Kingdom	2010-12-01	12
536367	84969	BOX OF 6 ASSORTED...	6	12/1/2010 8:34	4.25	13047	United Kingdom	2010-12-01	12
536367	22623	BOX OF VINTAGE JI...	3	12/1/2010 8:34	4.95	13047	United Kingdom	2010-12-01	12
536367	22622	BOX OF VINTAGE AL...	2	12/1/2010 8:34	9.95	13047	United Kingdom	2010-12-01	12
536367	21754	HOME BUILDING BLO...	3	12/1/2010 8:34	5.95	13047	United Kingdom	2010-12-01	12

Command took 1.22 seconds -- by alielsherifeng@gmail.com at 12/16/2022, 11:20:53 PM on Ali El-Sherif

Showing the monthly invoice aggregate summary:

```
val invoice_summary = sorted_data.groupBy("month").agg(
  sum("Quantity").as("total products sold"),
  avg("UnitPrice").as("average price"),
  countDistinct("CustomerID").as("total customers"),
  sum((sorted_data("Quantity")*sorted_data("UnitPrice"))).as("sales value")
)
display(invoice_summary)
```

▶ (3) Spark Jobs

▶  invoice_summary: org.apache.spark.sql.DataFrame = [month: integer, total products sold: long ... 3 more fields]

Table  +

	month	total products sold	average price	total customers	sales value
1	12	79062	3.8177434936908563	323	181847.24999999999

 Showing 1 row. | 2.89 seconds runtime

Command took 2.89 seconds -- by alielsherifeng@gmail.com at 12/16/2022, 11:21:32 PM on Ali El-Sherif

Question 3:

a)

Since disk size = 512, log size = 112, replication factor = 3

Available disk size per node = $512 - 112 = 400$ GB/node

Number of data nodes = $\frac{(300 \times 1024)}{400} \times 3 = 2304$

Total number of nodes = $2304 + 768 = 3072$

b)

Since log files = 640 GB, 4 containers per node

let the number of containers = number of partitions

so that available memory in HDFS per node = $64 - 14 = 50$ GB/node

number of nodes = $\frac{640}{50} = 12.8 \cong 13$ nodes.

number of partitions = $13 \times 4 = 52$ partitions

Question 4:

a)

$$\text{Number of sensors} = \frac{\text{Pipeline length}}{\text{maintaenance point distance}} = \frac{6000}{1.5} = 4000 \text{ sensors}$$

b)

$$\text{Number of total events} = \frac{4000}{30} = 133.333 \text{ events per second}$$

$$\text{Handled by each partition} = \frac{1500}{60} = 25 \text{ event per second}$$

$$\text{Number of partitions} = \frac{\text{Number of total events}}{\text{Handled by each partition}} = 5.333 \cong 6 \text{ partitions}$$

c)

According to the documentation [1], in basic or standard tier:

we will need 1 event hub with number of partitions = 32 (6 partitions needed)

Number of event hubs per namespace = 10, so 1 namespace is enough in this case

d)

Job query for leakage (<800 PSI):

```
SELECT SerialNumber, Longitude, Latitude, ReadingTime,  
PressureReading, Count(*) as CountofSensors  
FROM StreamData  
TIMESTAMP BY ReadingTime  
WHERE PressureReading < 800  
GROUP BY SerialNumber, SlidingWindow(Minute,5)  
HAVING CountofSensors >= 3
```

Job query for blockage (>1200 PSI):

```
SELECT SerialNumber, Longitude, Latitude, ReadingTime,  
PressureReading, Count(*) as CountofSensors  
FROM StreamData  
TIMESTAMP BY ReadingTime  
WHERE PressureReading >1200  
GROUP BY SerialNumber, SlidingWindow(Minute,5)  
HAVING CountofSensors >= 3
```

e)

Since there are 2 actions that can be taken, 2 consumer groups are needed and can be divided as follows:

- In case of obstruction or blockage downstream ($>1200\text{ PSI}$), the pump will shut off until pressure normalizes.
- In case of leakage ($< 800\text{ PSI}$), an alert will be triggered for the maintenance point associated with the sensor.

References

- [1] spelluru. (n.d.). *Quotas and limits - Azure Event Hubs - Azure Event Hubs*. Learn.microsoft.com.
<https://learn.microsoft.com/en-us/azure/event-hubs/event-hubs-quotas>